

---

# L5: Augmenting Data

Professor Daniel Yue  
Information Technology Management Area  
Georgia Tech, Scheller College of Business  
MGMT 7609 – Observational Studies of Information Systems

# This Week

L1  
Planning  
Research

L2  
Building  
Data  
Pipelines  
(Databases)

L2 Collecting Data (Scrapers)

L4 Exploring Data (Viz)

L5 Augmenting Data  
(LLMs)

L6 Merging Data  
(Crosswalks)

L7  
Student  
Presentations

---

# Motivation

---

## The Big Picture: LLMs in Research

- L2-L4 focused on using LLMs/AI tools as an assistant – boosting productivity in writing code, implementing our data pipelines and exploring the data.
- This is distinct from today's (and next week's) focus, using LLMs as a tool within a research design.

---

# Two Ways LLMs Function as Research Tools

## LLMs as Data Processors (Prompt-to-Response)

- The LLM directly provides the output we want.
- We give it raw text and a prompt.
- It gives us a label, a summary, or structured data.

*Focus for today – structure extraction of information*

## LLMs as Feature Generators (Embeddings)

- E.g. BERT, Sentence Transformers
- Turns text into dense vectors (numbers).
- These vectors become features used in downstream tasks (e.g., matching, regression).

*Focus for next week – application to building crosswalks*



## Goals for this lecture:

Introduction to using LLMs for text-based research tasks

### Processing Text with LLMs

LLMs as tools in research designs

“Text as Data” (Gentzkow, Kelly, Taddy 2019)

“LLMs to Annotate Data”  
(Carlson and Burbano 2025)

### Extracting Structured Data with LLMs

Practical Application

Extracting data from HF  
Model Cards

---

# Processing Data with LLMs



# Data Augmentation

\*Not discussed today but important:  
Topic Modeling (discovering latent categories)

We are not generating new raw text; we are ***augmenting*** an existing dataset

- by applying judgment to that data (classification)
- by adding structure to unstructured data (structured extraction)

## Example 1: Classification

(known categories)

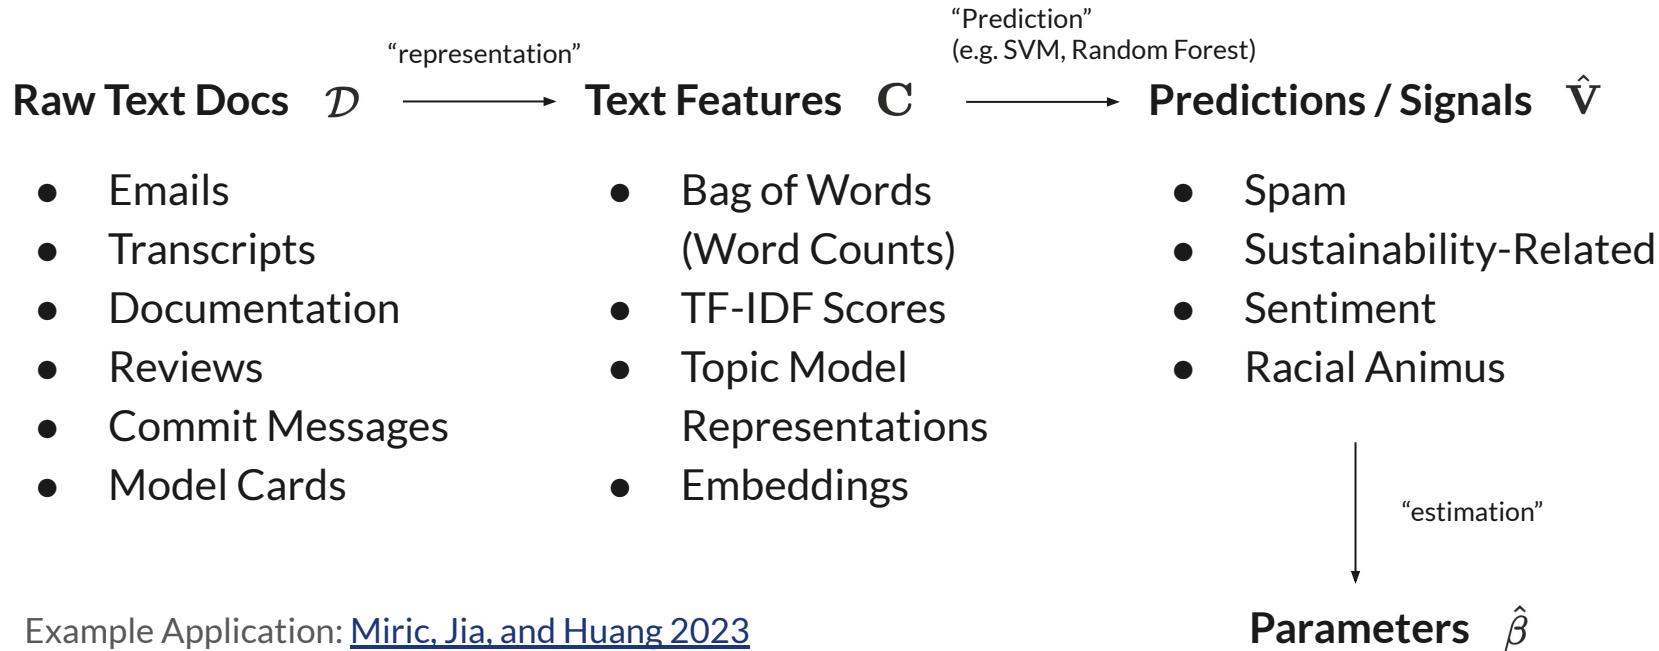
- News Article → Republican or Democrat
- Model Card → Did this model use a proprietary dataset?

## Example 2: Structured Extraction

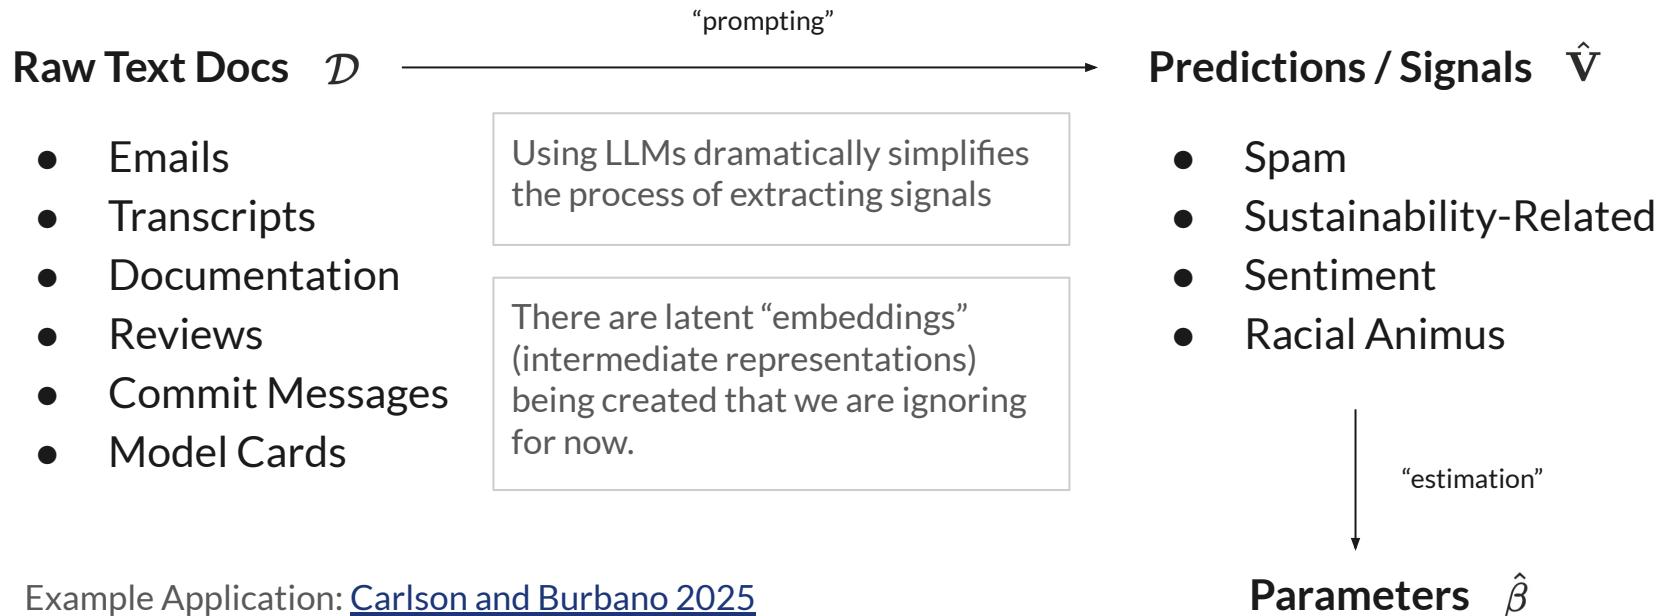
(unknown categories)

- News Article → List of all people mentioned in the article
- Model Card → Which evaluation metrics are reported?

# Text as Data (Gentzkow, Kelly, & Taddy 2019)



# Text as Data (Gentzkow, Kelly, & Taddy 2019) with LLMs!



---

# Helpful Framework: Carlson and Burbano 2025

1. **Method Selection.** What signal do you need? (Research-design driven)
2. **Model Selection.** Which model to use? Several considerations, and consider using multiple!
3. **Prompt Engineering.** Draft your prompt and define your target JSON output structure.
  - Manually label a small, random sample ( $n=100$ ) of texts (this is your "gold standard").
  - Run your LLM script over the *same* 100 texts.
  - Compare LLM output to human labels. Calculate accuracy. Analyze errors. Refine your prompt/schema.
4. **Cost and Scale Considerations**
  - Once validation is high (e.g., >95% agreement), run on your full dataset.
5. **Analytical Validation and Robustness.**
  - Check robustness of your parameters to prompting choices

Research stage	Principle(s)	Tradeoffs	Best practices
1. Method selection and integration ( <i>Do we use LLMs or another method?</i> )	Consider LLMs as part of a broader methodological toolkit; match method to research context and data characteristics	<ul style="list-style-type: none"><li>• Human Coding: High accuracy and nuance but expensive/slow/limited scale</li><li>• Keywords/Dictionary: Transparent and simple but rigid/context-blind</li><li>• Supervised ML: Scalable and consistent but requires training data/domain expertise</li><li>• LLMs: Flexible and fast but potentially unstable/opaque</li></ul>	<ul style="list-style-type: none"><li>• Begin with clear task specification and success criteria</li><li>• Use human coding for validation and complex edge cases</li><li>• Consider hybrid approaches</li><li>• Document relative performance across methods</li></ul>

Research stage	Principle(s)	Tradeoffs	Best practices
2. Model selection and stability ( <i>Which LLM should we be using?</i> )	Prioritize reproducibility and documentation over cutting-edge performance	<ul style="list-style-type: none"><li>• Newer models: Better performance but less tested/stable</li><li>• Older models: More stable but potentially lower performance</li><li>• Open vs. closed models: Access/cost vs. performance</li><li>• Reasoning models: Optimized for complex reasoning vs. uncertain effectiveness for basic annotation tasks</li></ul>	<ul style="list-style-type: none"><li>• Document model version and access method</li><li>• Test multiple models as part of sensitivity analysis</li><li>• Archive model weights (if open-source) or maintain access to older versions if possible</li><li>• Plan for replication needs</li><li>• Consider data security if annotating sensitive data (check policies around usage and storage)</li></ul>

---

## **Key Technical Skill #1: Calling LLMs through an API**

1.

Research stage	Principle(s)	Tradeoffs	Best practices
3. Prompt engineering ( <i>How do we design and validate our prompts?</i> )	Structure prompts systematically and transparently; document thoroughly	<ul style="list-style-type: none"><li>Simple prompts: More stable but potentially lower accuracy</li><li>Complex prompts: Better performance but harder to replicate</li><li>Generic vs. task-specific prompting: Generalizability vs. accuracy</li></ul>	<ul style="list-style-type: none"><li>Use established frameworks (Chain of Thought, few-shot learning)</li><li>Generate systematic prompt variations for sensitivity testing</li><li>Document all prompt components</li><li>Test prompts on diverse examples</li></ul>

---

## **Key Issues: Correctness, Improvement, and Reliability**

1. How do we know the LLM output is correct?
2. How do we make the LLM output better?
3. How do we interpret / constrain the LLM output format?
  - Imposing JSON format
  - Imposing specific categories

---

## Correctness – Validating the Output

*How would you do this?*

- Labelled Examples! Need to manually create, can be time consuming.
  - For classification, the labels
  - For structured extraction, the true values
  - ~100 example is usually sufficient (rule of thumb)
- Best practice: create “training” set and “holdout” set to avoid overfitting
  - As if you were training a supervised ML model

---

# Improvement – Making It Better

*How would you do this?*

1. A few critical tricks:
  - System Prompt – provide context
  - Few-Shot Prompting (examples) – ***error analysis*** can help here
    - I cannot stress enough how important it is to look at the outputs.
  - Chain-of-thought / suggested reasoning path
  - (Potentially) Tools – enable LLM to seek additional context
2. Better Models!
  - Experiment with a few different models
  - Often, fine-tuned smaller models can outperform larger models  
(but requires sufficient labels)



## Reliability – Constraining the Output

*How would you do this?*

1. More Prompt Engineering Tricks!
  - Clear instructions
  - Few-Shot Prompting (examples) is extremely powerful
2. Structured Output
  - Modern LLM APIs (like OpenRouter, OpenAI) can force the output to conform to a specific JSON schema.
  - How it works (briefly): The API uses a contextual grammar to guide token generation, ensuring the output is valid JSON.

---

## **Key Technical Skill #2: Structuring the Output**

1.

Research stage	Principle(s)	Tradeoffs	Best practices
4. Cost and scale considerations ( <i>How do we optimize resources while maintaining rigor?</i> )	Balance resource usage with statistical power and robustness needs	<ul style="list-style-type: none"><li>• Sample size vs. token usage</li><li>• Prompt complexity vs. processing cost</li><li>• Depth vs. breadth of sensitivity testing</li><li>• Model size vs. performance gains</li></ul>	<ul style="list-style-type: none"><li>• Plan sensitivity analysis within resource constraints</li><li>• Document cost structure and optimization decisions</li><li>• Consider computational constraints for replication</li></ul>

---

## **Key Technical Skill #3: Estimating the Cost**

1.

Research stage	Principle(s)	Tradeoffs	Best practices
5. Validation and robustness assessment ( <i>How reliable and stable are our findings?</i> )	Conduct systematic evaluation of result validity and stability	<ul style="list-style-type: none"><li>• Depth vs. breadth of validation testing</li><li>• Cost of comprehensive testing vs. confidence in results</li><li>• Simple vs. complex sensitivity metrics</li></ul>	<ul style="list-style-type: none"><li>• If feasible, maintain gold-standard expert-coded subset</li><li>• Generate systematic prompt variations</li><li>• Test impact on downstream analyses</li><li>• Provide bounded estimates for key findings</li><li>• Document null results and validation protocols</li></ul>

This part is the wild west in our field right now. Lots of ideas, but unclear what referees may ask for.  
I would not over-engineer on this aspect yet.

---

# Review of Carlson and Burbano 2025

1. **Method Selection.** What signal do you need? (Research-design driven)
2. **Model Selection.** Which model to use? Several considerations, and consider using multiple!
3. **Prompt Engineering.** Draft your prompt and define your target JSON output structure.
  - Manually label a small, random sample ( $n=100$ ) of texts (this is your "gold standard").
  - Run your LLM script over the *same* 100 texts.
  - Compare LLM output to human labels. Calculate accuracy. Analyze errors. Refine your prompt/schema.
4. **Cost and Scale Considerations**
  - Once validation is high, run on your full dataset.
5. **Analytical Validation and Robustness.**
  - Check robustness of your parameters to prompting choices

---

**15 minute break!**

---

# Extracting Structured Data with LLMs

Extract data from the Model Cards on Hugging Face to characterize the nature of fine-tuning at scale.

[librarian-bots/model\\_cards\\_with\\_metadata](#)



---

## Assignment

Add one or more LLM-generated variables to your dataset.  
Document your validation process and create visualizations  
comparing automated labels to manual validation samples.

# Next Week

L1  
Planning  
Research

L2  
Building  
Data  
Pipelines  
(Databases)

L2 Collecting Data (Scrapers)

L4 Exploring Data (Viz)

L5 Augmenting Data  
(LLMs)

L6 Merging Data  
(Crosswalks)

L7  
Student  
Presentations

---

## Class Prep

- Revealing the revealed preferences of public firm CEOs and top executives: A new database from credit card spending – Raffiee et al 2022
  - Section 2.2 and 2.3
- Deep Learning for Economists – Dell 2024
  - Section VIII.1-4 (Embedding Models)
- MTEB: Massive Text Embedding Benchmark – Muennighoff et al 2023
  - Short, influential benchmark – worth reading!